

✓ Softmax Function:

```

0s [1] import numpy as np

def softmax(z):
    """
    Compute the softmax probabilities for a given input matrix.

    Parameters:
    z (numpy.ndarray): Logits (raw scores) of shape (m, n), where
        - m is the number of samples.
        - n is the number of classes.

    Returns:
    numpy.ndarray: Softmax probability matrix of shape (m, n), where
        each row sums to 1 and represents the probability
        distribution over classes.

    Notes:
    - The input to softmax is typically computed as:  $z = XW + b$ .
    - Uses numerical stabilization by subtracting the max value per row.
    """

    z_shifted = z - np.max(z, axis=1, keepdims=True)
    exp_z = np.exp(z_shifted)
    return exp_z / np.sum(exp_z, axis=1, keepdims=True)

```

✓ Softmax Test Case:

This test case checks that each row in the resulting softmax probabilities sums to 1, which is the fundamental property of softmax.

```

0s [2] # Example test case
z_test = np.array([[2.0, 1.0, 0.1], [1.0, 1.0, 1.0]])
softmax_output = softmax(z_test)
print(softmax_output)

# Verify if the sum of probabilities for each row is 1 using assert
row_sums = np.sum(softmax_output, axis=1)
print(row_sums)

# Assert that the sum of each row is 1
assert np.allclose(row_sums, 1), f"Test failed: Row sums are {row_sums}"

print("Softmax function passed the test case!")

```

```

[[0.65900114 0.24243297 0.09856589]
 [0.33333333 0.33333333 0.33333333]]
[1. 1.]
Softmax function passed the test case!

```

✓ Prediction Function:

```
✓ [3] def predict_softmax(X, W, b):  
    """  
    Predict the class labels for a set of samples using the trained softmax model.  
  
    Parameters:  
    X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of samples and d is the number of features.  
    W (numpy.ndarray): Weight matrix of shape (d, c), where c is the number of classes.  
    b (numpy.ndarray): Bias vector of shape (c,).  
  
    Returns:  
    numpy.ndarray: Predicted class labels of shape (n,), where each value is the index of the predicted class.  
    """  
    z = np.dot(X, W) + b # Compute the scores (logits)  
    y_pred = softmax(z) # Get the probabilities using the softmax function  
  
    # Assign the class with the highest probability  
    predicted_classes = np.argmax(y_pred, axis=1)  
  
    return predicted_classes
```

✓ Test Function for Prediction Function:

The test function ensures that the predicted class labels have the same number of elements as the input samples, verifying that the model produces a valid output shape.

```
✓ [4] # Define test case  
X_test = np.array([[0.2, 0.8], [0.5, 0.5], [0.9, 0.1]]) # Feature matrix (3 samples, 2 features)  
W_test = np.array([[0.4, 0.2, 0.1], [0.3, 0.7, 0.5]]) # Weights (2 features, 3 classes)  
b_test = np.array([0.1, 0.2, 0.3]) # Bias (3 classes)  
  
# Expected Output:  
# The function should return an array with class labels (0, 1, or 2)  
  
y_pred_test = predict_softmax(X_test, W_test, b_test)  
  
# Validate output shape  
assert y_pred_test.shape == (3,), f"Test failed: Expected shape (3,), got {y_pred_test.shape}"  
  
# Print the predicted labels  
print("Predicted class labels:", y_pred_test)
```

➡ Predicted class labels: [1 1 0]

▼ Loss Function:

```
[5] def loss_softmax(y_pred, y):  
    """  
    Compute the cross-entropy loss for a single sample.  
  
    Parameters:  
    y_pred (numpy.ndarray): Predicted probabilities of shape (c,) for a single sample,  
                            where c is the number of classes.  
    y (numpy.ndarray): True labels (one-hot encoded) of shape (c,), where c is the number of classes.  
  
    Returns:  
    float: Cross-entropy loss for the given sample.  
    """  
  
    epsilon = 1e-12 # To avoid log(0)  
    y_pred = np.clip(y_pred, epsilon, 1.0 - epsilon) # Prevent log(0) by clipping values  
    n = y.shape[0] # Number of samples  
    loss = -np.sum(y * np.log(y_pred)) / n  
    return loss
```

▼ Test case for Loss Function:

This test case Compares loss for correct vs. incorrect predictions.

- Expects low loss for correct predictions.
- Expects high loss for incorrect predictions.

```
[6] import numpy as np  
  
# Define correct predictions (low loss scenario)  
y_true_correct = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]]) # True one-hot labels  
y_pred_correct = np.array([[0.9, 0.05, 0.05],  
                           [0.1, 0.85, 0.05],  
                           [0.05, 0.1, 0.85]]) # High confidence in the correct class  
  
# Define incorrect predictions (high loss scenario)  
y_pred_incorrect = np.array([[0.05, 0.05, 0.9], # Highly confident in the wrong class  
                             [0.1, 0.05, 0.85],  
                             [0.85, 0.1, 0.05]])  
  
# Compute loss for both cases  
loss_correct = loss_softmax(y_pred_correct, y_true_correct)  
loss_incorrect = loss_softmax(y_pred_incorrect, y_true_correct)  
  
# Validate that incorrect predictions lead to a higher loss  
assert loss_correct < loss_incorrect, f"Test failed: Expected loss_correct < loss_incorrect, but got {loss_correct:.4f} >= {loss_incorrect:.4f}"  
  
# Print results  
print(f"Cross-Entropy Loss (Correct Predictions): {loss_correct:.4f}")  
print(f"Cross-Entropy Loss (Incorrect Predictions): {loss_incorrect:.4f}")
```

```
📄 Cross-Entropy Loss (Correct Predictions): 0.1435  
Cross-Entropy Loss (Incorrect Predictions): 2.9957
```

Cost Function:

```
def cost_softmax(X, y, W, b):  
    """  
    Compute the average softmax regression cost (cross-entropy loss) over all samples.  
  
    Parameters:  
    X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of samples and d is the number of features.  
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c), where n is the number of samples and c is the number of classes.  
    W (numpy.ndarray): Weight matrix of shape (d, c).  
    b (numpy.ndarray): Bias vector of shape (c,).  
  
    Returns:  
    float: Average softmax cost (cross-entropy loss) over all samples.  
    """  
  
    n = X.shape[0] # Number of samples  
    z = np.dot(X, W) + b  
    y_pred = softmax(z)  
    cost = loss_softmax(y_pred, y)  
    return cost
```

Test Case for Cost Function:

The test case assures that the cost for the incorrect prediction should be higher than for the correct prediction, confirming that the cost function behaves as expected.

```
import numpy as np  
  
# Example 1: Correct Prediction (Closer predictions)  
X_correct = np.array([[1.0, 0.0], [0.0, 1.0]]) # Feature matrix for correct predictions  
y_correct = np.array([[1, 0], [0, 1]]) # True labels (one-hot encoded, matching predictions)  
W_correct = np.array([[5.0, -2.0], [-3.0, 5.0]]) # Weights for correct prediction  
b_correct = np.array([0.1, 0.1]) # Bias for correct prediction  
  
# Example 2: Incorrect Prediction (Far off predictions)  
X_incorrect = np.array([[0.1, 0.9], [0.8, 0.2]]) # Feature matrix for incorrect predictions  
y_incorrect = np.array([[1, 0], [0, 1]]) # True labels (one-hot encoded, incorrect predictions)  
W_incorrect = np.array([[0.1, 2.0], [1.5, 0.3]]) # Weights for incorrect prediction  
b_incorrect = np.array([0.5, 0.6]) # Bias for incorrect prediction  
  
# Compute cost for correct predictions  
cost_correct = cost_softmax(X_correct, y_correct, W_correct, b_correct)  
  
# Compute cost for incorrect predictions  
cost_incorrect = cost_softmax(X_incorrect, y_incorrect, W_incorrect, b_incorrect)  
  
# Check if the cost for incorrect predictions is greater than for correct predictions  
assert cost_incorrect > cost_correct, f"Test failed: Incorrect cost {cost_incorrect} is not greater than correct cost {cost_correct}"  
  
# Print the costs for verification  
print("Cost for correct prediction:", cost_correct)  
print("Cost for incorrect prediction:", cost_incorrect)  
  
print("Test passed!")  
  
Cost for correct prediction: 0.0006234364133349324  
Cost for incorrect prediction: 0.29930861359446115  
Test passed!
```

```

[9] def compute_gradient_softmax(X, y, W, b):
    """
    Compute the gradients of the cost function with respect to weights and biases.

    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d).
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).
    W (numpy.ndarray): Weight matrix of shape (d, c).
    b (numpy.ndarray): Bias vector of shape (c,).

    Returns:
    tuple: Gradients with respect to weights (d, c) and biases (c,).
    """
    n, d = X.shape
    z = np.dot(X, W) + b
    y_pred = softmax(z)

    grad_W = np.dot(X.T, (y_pred - y)) / n # Gradient with respect to weights
    grad_b = np.sum(y_pred - y, axis=0) / n # Gradient with respect to biases

    return grad_W, grad_b

```

▼ Test case for compute_gradient function:

The test checks if the gradients from the function are close enough to the manually computed gradients using `np.allclose`, which accounts for potential floating-point discrepancies.

```

import numpy as np

# Define a simple feature matrix and true labels
X_test = np.array([[0.2, 0.8], [0.5, 0.5], [0.9, 0.1]]) # Feature matrix (3 samples, 2 features)
y_test = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]]) # True labels (one-hot encoded, 3 classes)

# Define weight matrix and bias vector
W_test = np.array([[0.4, 0.2, 0.1], [0.3, 0.7, 0.5]]) # Weights (2 features, 3 classes)
b_test = np.array([0.1, 0.2, 0.3]) # Bias (3 classes)

# Compute the gradients using the function
grad_W, grad_b = compute_gradient_softmax(X_test, y_test, W_test, b_test)

# Manually compute the predicted probabilities (using softmax function)
z_test = np.dot(X_test, W_test) + b_test
y_pred_test = softmax(z_test)

# Compute the manually computed gradients
grad_W_manual = np.dot(X_test.T, (y_pred_test - y_test)) / X_test.shape[0]
grad_b_manual = np.sum(y_pred_test - y_test, axis=0) / X_test.shape[0]

# Assert that the gradients computed by the function match the manually computed gradients
assert np.allclose(grad_W, grad_W_manual), f"Test failed: Gradients w.r.t. W are not equal.\nExpected: {grad_W_manual}\nGot: {grad_W}"
assert np.allclose(grad_b, grad_b_manual), f"Test failed: Gradients w.r.t. b are not equal.\nExpected: {grad_b_manual}\nGot: {grad_b}"

# Print the gradients for verification
print("Gradient w.r.t. W:", grad_W)
print("Gradient w.r.t. b:", grad_b)

print("Test passed!")

```

Gradient w.r.t. W: [[0.1031051 0.01805685 -0.12116196]
 [-0.13600547 0.00679023 0.12921524]]
 Gradient w.r.t. b: [-0.03290036 0.02484708 0.00805328]
 Test passed!

✓ Implementing Gradient Descent:

```
[11] def gradient_descent_softmax(X, y, W, b, alpha, n_iter, show_cost=False):  
    """  
    Perform gradient descent to optimize the weights and biases.  
  
    Parameters:  
    X (numpy.ndarray): Feature matrix of shape (n, d).  
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).  
    W (numpy.ndarray): Weight matrix of shape (d, c).  
    b (numpy.ndarray): Bias vector of shape (c,).  
    alpha (float): Learning rate.  
    n_iter (int): Number of iterations.  
    show_cost (bool): Whether to display the cost at intervals.  
  
    Returns:  
    tuple: Optimized weights, biases, and cost history.  
    """  
    cost_history = []  
  
    for i in range(n_iter):  
        # Compute gradients  
        grad_W, grad_b = compute_gradient_softmax(X, y, W, b)  
  
        # Update weights and biases using the gradients  
        W -= alpha * grad_W  
        b -= alpha * grad_b  
  
        # Compute and store cost  
        cost = cost_softmax(X, y, W, b)  
        cost_history.append(cost)  
  
        # Print cost at regular intervals  
        if show_cost and (i % 100 == 0 or i == n_iter - 1):  
            print(f"Iteration {i}: Cost = {cost:.6f}")  
  
    return W, b, cost_history
```

✓ Preparing Dataset:

```
✓ [12] import pandas as pd
1s      import numpy as np
      import matplotlib.pyplot as plt
      from sklearn.model_selection import train_test_split

      def load_and_prepare_mnist(csv_file, test_size=0.2, random_state=42):
          """
          Reads the MNIST CSV file, splits data into train/test sets, and plots one image per class.

          Arguments:
          csv_file (str)      : Path to the CSV file containing MNIST data.
          test_size (float)   : Proportion of the data to use as the test set (default: 0.2).
          random_state (int)  : Random seed for reproducibility (default: 42).

          Returns:
          X_train, X_test, y_train, y_test : Split dataset.
          """

          # Load dataset
          df = pd.read_csv(csv_file)

          # Separate labels and features
          y = df.iloc[:, 0].values # First column is the label
          X = df.iloc[:, 1:].values # Remaining columns are pixel values

          # Normalize pixel values (optional but recommended)
          X = X / 255.0 # Scale values between 0 and 1

          # Split data into train and test sets
          X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size, random_state=random_state)

          # Plot one sample image per class
          plot_sample_images(X, y)

          return X_train, X_test, y_train, y_test

      def plot_sample_images(X, y):
          """
          Plots one sample image for each digit class (0-9).

          Arguments:
```

```
          Arguments:
          X (np.ndarray): Feature matrix containing pixel values.
          y (np.ndarray): Labels corresponding to images.
          """

          plt.figure(figsize=(10, 4))
          unique_classes = np.unique(y) # Get unique class labels

          for i, digit in enumerate(unique_classes):
              index = np.where(y == digit)[0][0] # Find first occurrence of the class
              image = X[index].reshape(28, 28) # Reshape 1D array to 28x28


              plt.subplot(2, 5, i + 1)
              plt.imshow(image, cmap='gray')
              plt.title(f"Digit: {digit}")
              plt.axis('off')

          plt.tight_layout()
          plt.show()
```

```
✓ [17] from google.colab import drive
1s      drive.mount('/content/drive')
```

➡ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
✓ 10s [18] csv_file_path = "/content/drive/MyDrive/Artificial Intelligence and Machine Learning/mnist_dataset.csv" # Path to saved dataset
X_train, X_test, y_train, y_test = load_and_prepare_mnist(csv_file_path)
```



✓ A Quick debugging Step:

```
✓ 0s [19] # Assert that X and y have matching lengths
assert len(X_train) == len(y_train), f"Error: X and y have different lengths! X={len(X_train)}, y={len(y_train)}"
print("Move forward: Dimension of Feature Matrix X and label vector y matched.")
```

↻ Move forward: Dimension of Feature Matrix X and label vector y matched.

▼ Train the Model:

```
✓ [20] print(f"Training data shape: {X_train.shape}")  
0s      print(f"Test data shape: {X_test.shape}")
```


```
↗ Training data shape: (48000, 784)  
Test data shape: (12000, 784)
```

```
✓ [21] from sklearn.preprocessing import OneHotEncoder  
7m  
  
# Check if y_train is one-hot encoded  
if len(y_train.shape) == 1:  
    encoder = OneHotEncoder(sparse_output=False) # Use sparse_output=False for newer versions of sklearn  
    y_train = encoder.fit_transform(y_train.reshape(-1, 1)) # One-hot encode labels  
    y_test = encoder.transform(y_test.reshape(-1, 1)) # One-hot encode test labels  
  
# Now y_train is one-hot encoded, and we can proceed to use it  
d = X_train.shape[1] # Number of features (columns in X_train)  
c = y_train.shape[1] # Number of classes (columns in y_train after one-hot encoding)  
  
# Initialize weights with small random values and biases with zeros  
W = np.random.randn(d, c) * 0.01 # Small random weights initialized  
b = np.zeros(c) # Bias initialized to 0  
  
# Set hyperparameters for gradient descent  
alpha = 0.1 # Learning rate  
n_iter = 1000 # Number of iterations to run gradient descent  
  
# Train the model using gradient descent  
W_opt, b_opt, cost_history = gradient_descent_softmax(X_train, y_train, W, b, alpha, n_iter, show_cost=True)  
  
# Plot the cost history to visualize the convergence  
plt.plot(cost_history)  
plt.title('Cost Function vs. Iterations')  
plt.xlabel('Iterations')  
plt.ylabel('Cost')  
plt.grid(True)  
plt.show()
```



Iteration 0: Cost = 2.183171
Iteration 100: Cost = 0.606677
Iteration 200: Cost = 0.489237
Iteration 300: Cost = 0.440693
Iteration 400: Cost = 0.412645
Iteration 500: Cost = 0.393799
Iteration 600: Cost = 0.379994
Iteration 700: Cost = 0.369299
Iteration 800: Cost = 0.360685
Iteration 900: Cost = 0.353544
Iteration 999: Cost = 0.347549

Cost Function vs. Iterations



✓ Evaluating the Model:

✓
0s

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score

# Evaluation Function
def evaluate_classification(y_true, y_pred):
    """
    Evaluate classification performance using confusion matrix, precision, recall, and F1-score.

    Parameters:
    y_true (numpy.ndarray): True labels
    y_pred (numpy.ndarray): Predicted labels

    Returns:
    tuple: Confusion matrix, precision, recall, F1 score
    """
    # Compute confusion matrix
    cm = confusion_matrix(y_true, y_pred)

    # Compute precision, recall, and F1-score
    precision = precision_score(y_true, y_pred, average='weighted')
    recall = recall_score(y_true, y_pred, average='weighted')
    f1 = f1_score(y_true, y_pred, average='weighted')

    return cm, precision, recall, f1
```

```

# Predict on the test set
y_pred_test = predict_softmax(X_test, W_opt, b_opt)

# Evaluate accuracy
y_test_labels = np.argmax(y_test, axis=1) # True labels in numeric form

# Evaluate the model
cm, precision, recall, f1 = evaluate_classification(y_test_labels, y_pred_test)

# Print the evaluation metrics
print("\nConfusion Matrix:")
print(cm)
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1-Score: {f1:.2f}")

# Visualizing the Confusion Matrix
fig, ax = plt.subplots(figsize=(12, 12))
cax = ax.imshow(cm, cmap='Blues') # Use a color map for better visualization

# Dynamic number of classes
num_classes = cm.shape[0]
ax.set_xticks(range(num_classes))
ax.set_yticks(range(num_classes))
ax.set_xticklabels([f'Predicted {i}' for i in range(num_classes)])
ax.set_yticklabels([f'Actual {i}' for i in range(num_classes)])

# Add labels to each cell in the confusion matrix
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(j, i, cm[i, j], ha='center', va='center', color='white' if cm[i, j] > np.max(cm) / 2 else 'black')

# Add grid lines and axis labels
ax.grid(False)
plt.title('Confusion Matrix', fontsize=14)
plt.xlabel('Predicted Label', fontsize=12)
plt.ylabel('Actual Label', fontsize=12)

# Adjust layout
plt.tight_layout()
plt.colorbar(cax)
plt.show()


```



```

Confusion Matrix:
[[1127   0   5   2   3  11   9   2  13   3]
 [  0 1276   7  11   1   5   1   4  16   1]
 [  1  15 1027  16  18   5  27  24  34   7]
 [  8   5  34 1051   1  53   9   8  29  21]
 [  1   5   7   1 1092   0  10   4   3  53]
 [ 22  14  12  43  13  920  14   7  46  13]
 [  6   2  10   1  11  15 1120   2  10   0]
 [  7  26  24   4  16   2   0 1184   7  29]
 [  9  27  13  34   9  33  13   6 1001  15]
 [  8   6  10  18  42   9   0   39  10 1052]]
Precision: 0.90
Recall: 0.90
F1-Score: 0.90

```



```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification, make_circles
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression

# Set random seed for reproducibility
np.random.seed(42)

# Generate linearly separable dataset
X_linear_separable, y_linear_separable = make_classification(n_samples=200, n_features=2,
n_informative=2,

n_redundant=0, n_clusters_per_class=1,
random_state=42)

# Split the data into training and testing sets
X_train_linear, X_test_linear, y_train_linear, y_test_linear = train_test_split(
X_linear_separable, y_linear_separable, test_size=0.2, random_state=42
)

# Train logistic regression model on linearly separable data
logistic_model_linear_separable = LogisticRegression()
logistic_model_linear_separable.fit(X_train_linear, y_train_linear)

# Generate non-linearly separable dataset (circles)
X_non_linear_separable, y_non_linear_separable = make_circles(n_samples=200, noise=0.1, factor=0.5,
random_state=42)

# Split the data into training and testing sets
X_train_non_linear, X_test_non_linear, y_train_non_linear, y_test_non_linear = train_test_split(
X_non_linear_separable, y_non_linear_separable, test_size=0.2, random_state=42
)

# Train logistic regression model on non-linearly separable data
logistic_model_non_linear_separable = LogisticRegression()
logistic_model_non_linear_separable.fit(X_train_non_linear, y_train_non_linear)

# Plot decision boundaries for linearly and non-linearly separable data
def plot_decision_boundary(ax, model, X, y, title):
    h = .02 # step size in the mesh
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    ax.contourf(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)
    ax.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=plt.cm.Paired)

```

11

